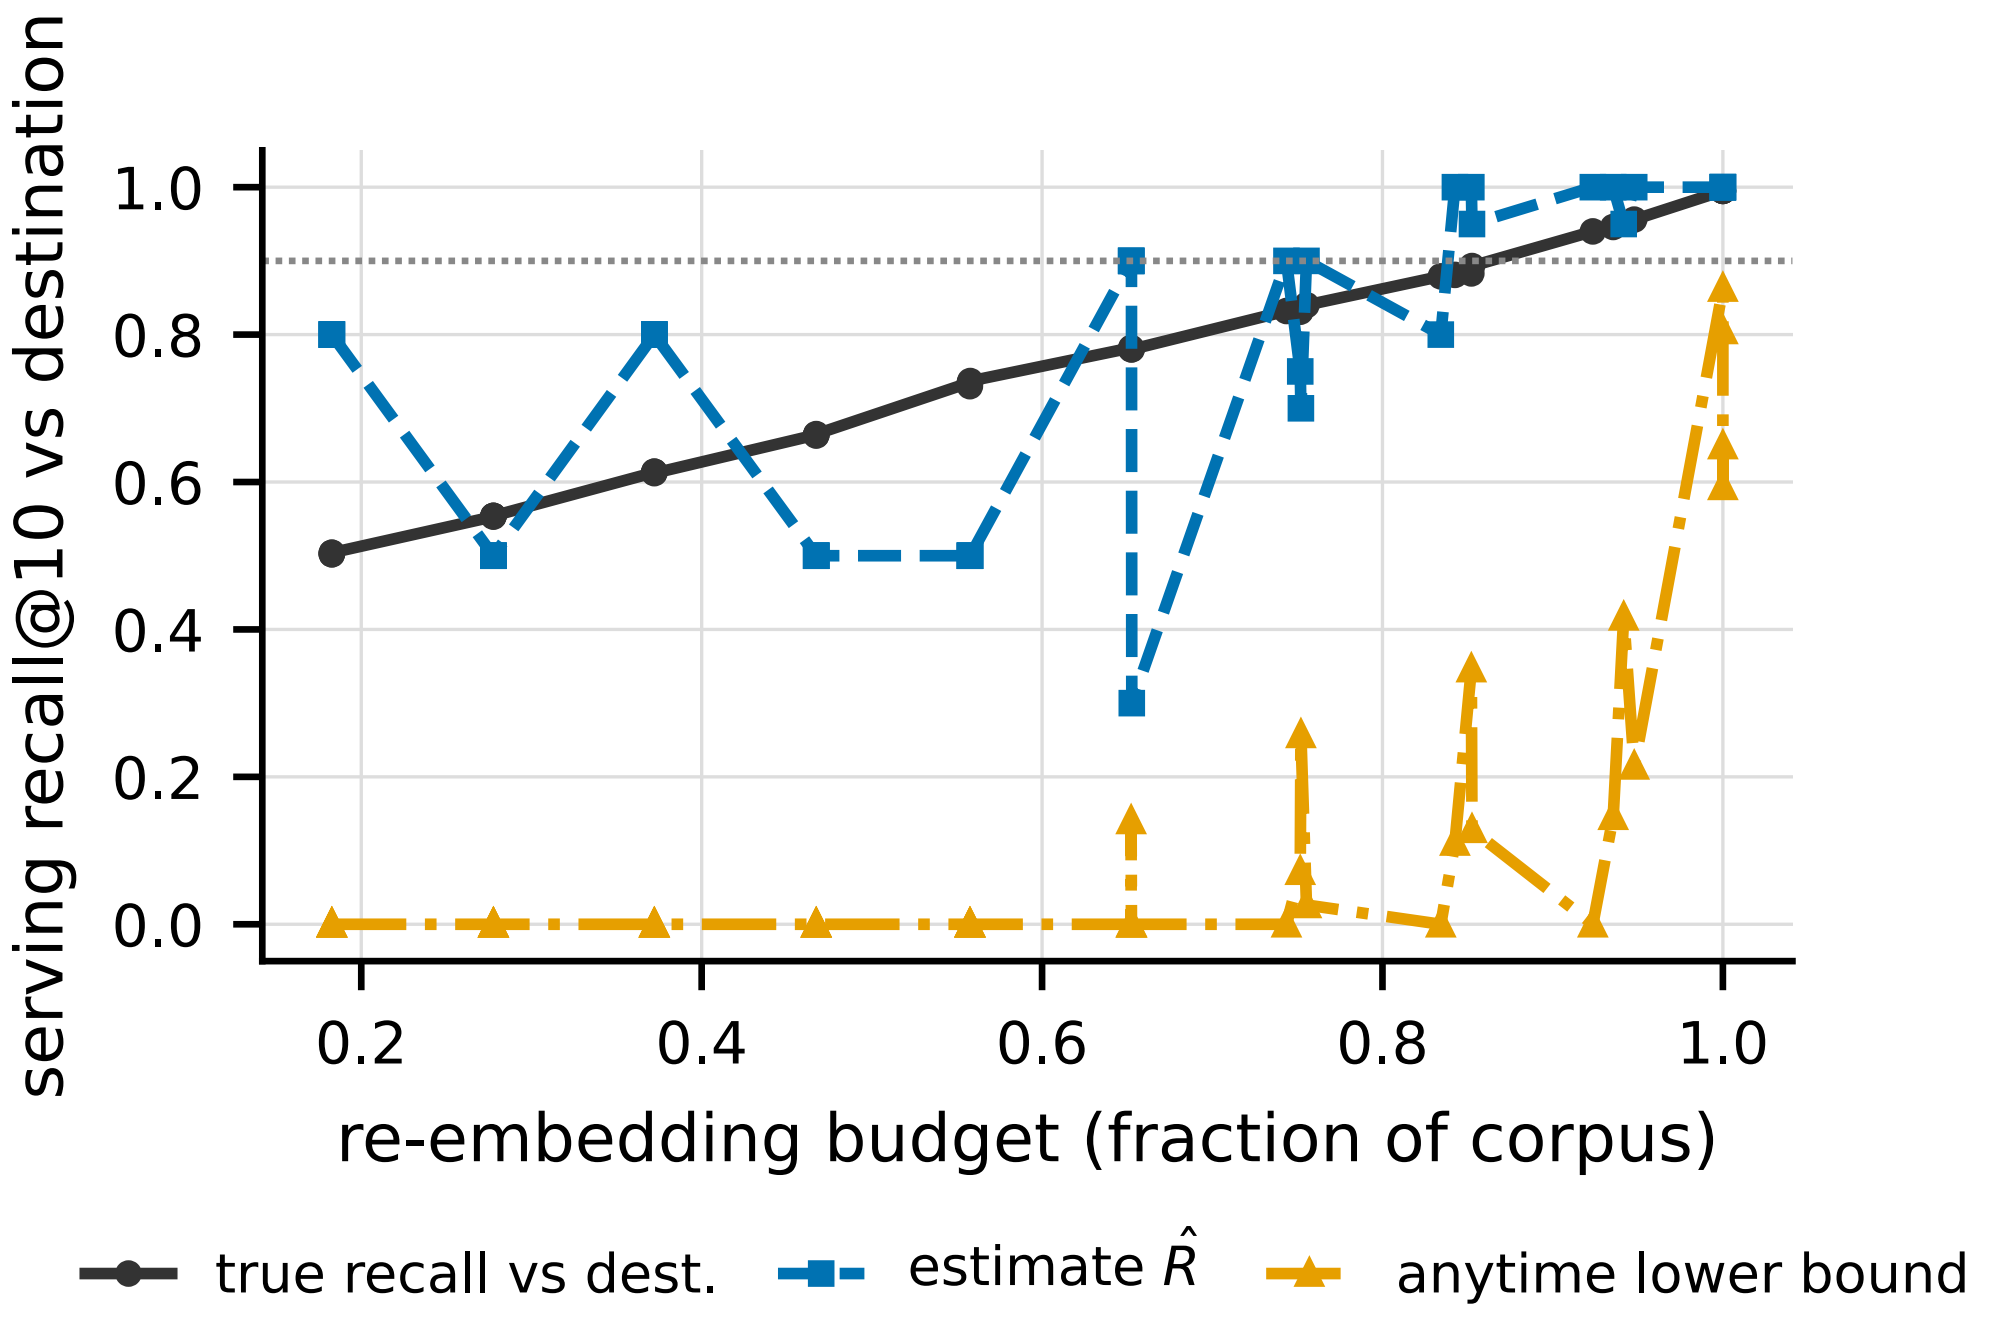
serving recall@10 vs destination
1.0
0.8
0.6
0.4
0.2
0.0
0.2
0.4
0.6
0.8
1.0
re-embedding budget (fraction of corpus)
true recall vs dest.
estimate $\hat{R}$
anytime lower bound